

BIG DATA (CS5620)

Introduction to MapReduce

Instructor: Dr. Yui Man Lui

Why MapReduce

Big Data

- Big data is everywhere
 - There are 2,161,530,000,000 searches in 2013
 - 92,100,000 pages mentioning Albert Einstein
 - Only 38,400 pages mentioning Yui Man Lui
- Gather as much data as you need and run machine learning / data mining / analytics algorithms to generate insights
- Challenges (3Vs)
 - Big Volume – the quantity of generated and stored data
 - Big Velocity – the speed at which the data is generated and processed
 - Big variety – the type and nature of the data

Big Data

- Need a fast way to process large amount of data
- MapReduce provides a framework for parallel computations using a large number of nodes
- MapReduce produces a scale-out, distributed, and fault-tolerant solution
- Hadoop / GFS is designed to use MapReduce as the primary method of interaction

What is MapReduce

MapReduce

- MapReduce is inspired by the concept of map and reduced in *Lisp*
 - A map function takes parameters as a function and a set of values
 - `(map 'length '(() (a) (ab) (abc))) → (0 1 2 3)`
 - The reduce function is a binary function and a set of values as parameters. It combines all values together using the given function
 - `(reduce #' + '(0 1 2 3)) → 6`

MapReduce

- MapReduce is a programming model
- MapReduce provides an abstraction to perform simple computations while hiding the details of parallelization, data distribution, load balancing, and fault tolerance
- Developed in Google as a mechanism for processing large amounts of data
 - Distributed data across thousands of machines
 - Same computations are performed on each CPU with different subsets of data

MapReduce

- The world has changed



Google Stanford Hardware (1998)



Google Data Center, Iowa

Need to think parallel

How MapReduce Works

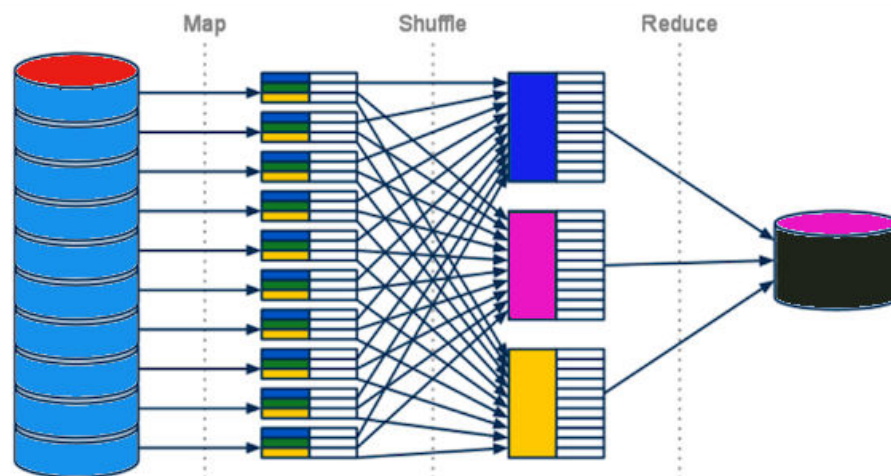
Mapper

- Mapper maps input (key, value) pairs to a set of intermediate (key, value) pairs
 - Maps are individual tasks that transform input data to intermediate records
 - A given input may map to zero or many output pairs

Reducer

- Reducer reduces a set of intermediate values which share a key (hash key) to a smaller set of values
- Reducer has three primary phases – Shuffle, Sort, Reduce
 - Shuffle – Take similar data and group them together
 - Sort – Groups the data to the reducer by keys
 - Reduce – Aggregate the data and summarize it in some way

Schematic of MapReduce Computation

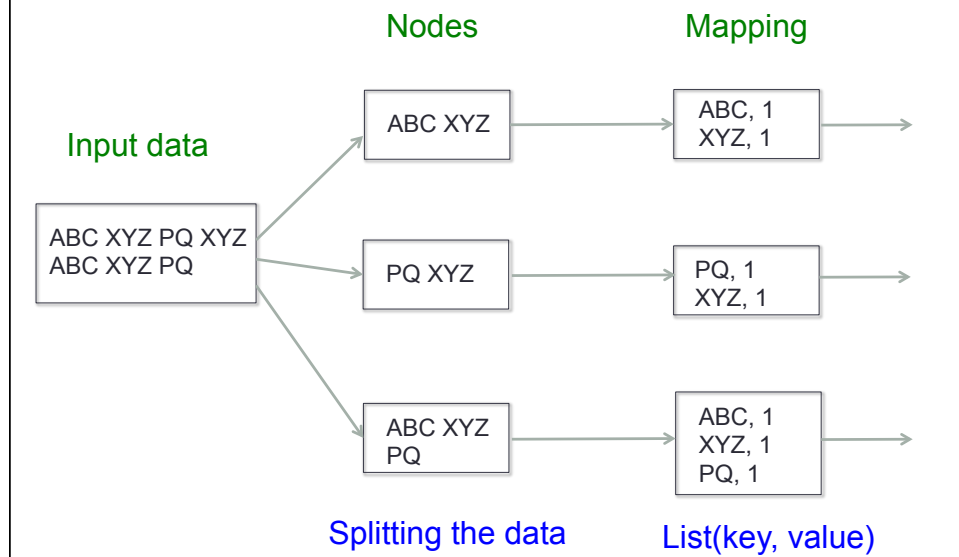


Schematic of MapReduce Computation

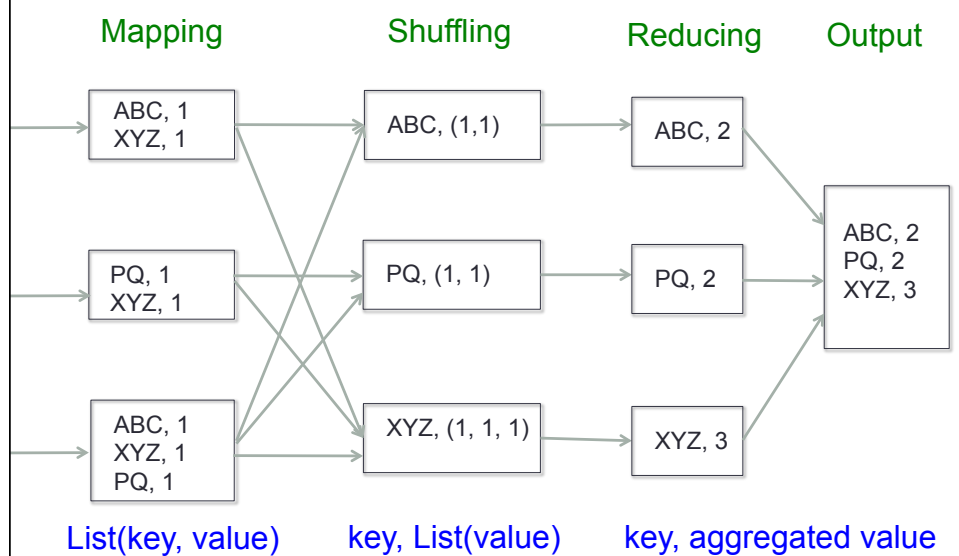
- Split phase – partitions the data across different mappers
 - Think they are different machines
- Each mapper executes *user defined* map code on the partitions in parallel
- Data is shuffled such that there is one reducer per key
 - Think they are different machines
- Each reducer executes the *user defined* reduce code in parallel

Word Count Example

MapReduce – Word Count Example



MapReduce – Word Count Example



Summary

- Input data distributed to nodes (servers)
- Each map works on a split data
- Data exchange between nodes in a “shuffle” process
- Intermediate data of the same key goes to the same reducer
- Reducer summarizes the values
- Reducer output is stored

- Need to provide a functional abstraction for mapper and reducer

Python - Word Count Example

MapReduce – Word Count Example

```
from multiprocessing import Pool

if __name__ == '__main__':
    text = loadtext('simple.txt')
    numProcessors = 3
    pool = Pool(processes=numProcessors)
    partitioned_text = chunks(text, 1 + len(text) / numProcessors)

    single_count_tuples = pool.map(Map, partitioned_text)
    token_to_tuples = Partition(single_count_tuples)
    term_frequencies = pool.map(Reduce, token_to_tuples.items())

    term_frequencies.sort(tuple_sort)

    for pair in term_frequencies:
        print('%20s: %5s' % (pair[0], pair[1]))
```

MapReduce – Word Count Example

```
def loadtext(path):
    word_list = []
    f = open(path, "r")
    for line in f:
        line = line.replace('.', ',')
        word_list.append(line)

    ## " ".join(["a", "b", "c"]).split() -> ['a', 'b', 'c']
    text = (" ".join(word_list)).split()
    return text

def chunks(text, n):
    data_block = []
    for i in xrange(0, len(text), n):
        data_block.append(text[i:i+n])

    return list(data_block)
```

MapReduce – Word Count Example

```
def Partition(lst):
    term_freq = {} # dictionary
    for sublist in lst:
        for w in sublist:
            ## if not a key in dictionary, then create a new key
            if (term_freq.get(w[0]) == None):
                term_freq[w[0]] = [w]
            else:
                ## otherwise append to the end
                term_freq[w[0]].append(w)

    return term_freq
```

MapReduce – Word Count Example

```
def Map(lst):
    results = []
    for word in lst:
        results.append((word, 1))

    return results

def Reduce(mapping):
    return (mapping[0], sum(pair[1] for pair in mapping[1]))

def tuple_sort (a, b):
    if (a[1] < b[1]):
        return 1
    elif (a[1] > b[1]):
        return -1
    else:
        return cmp(a[0], b[0])
```

MapReduce – Word Count Example

SINGLE COUNT TUPLES

```
[('ABC', 1), ('XYZ', 1), ('PQ', 1)], [('XYZ', 1), ('ABC', 1), ('XYZ', 1)], [('PQ', 1)]]
```

PARTITION DATA

```
{'XYZ': [('XYZ', 1), ('XYZ', 1), ('XYZ', 1)], 'ABC': [('ABC', 1), ('ABC', 1)], 'PQ':  
[('PQ', 1), ('PQ', 1)]}
```

TERM FREQUENCIES

```
[('XYZ', 3), ('ABC', 2), ('PQ', 2)]
```

```
XYZ: 3  
ABC: 2  
PQ: 2
```

Python – MRJOB Word Count Example

MapReduce – Word Count Example

```
from mrjob.job import MRJob
import re
WORD_RE = re.compile(r"[w']+")

class MRWordFreqCount(MRJob):
    def mapper(self, _, line):
        for word in WORD_RE.findall(line):
            ## we see this word one time
            yield (word, 1)

    def reducer(self, word, counts):
        yield (word, sum(counts))

if __name__ == '__main__':
    MRWordFreqCount.run()
```

MapReduce – Word Count Example

➤ python mr_wordcount-2.py simple.txt > output.txt

➤ cat output.txt

"ABC"	2
"PQ"	2
"XYZ"	3

Python – MRJOB Another Word Count Example

MapReduce – Word Count Example

```
from mrjob.job import MRJob

class MRWordFrequencyCount(MRJob):

    def mapper(self, _, line):
        yield "chars", len(line)
        yield "words", len(line.split())
        ## we see one line
        yield "lines", 1

    def reducer(self, key, values):
        yield key, sum(values)

if __name__ == '__main__':
    MRWordFrequencyCount.run()
```

MapReduce – Word Count Example

➤ `python mr_wordcount_example.py test.txt > output.txt`

➤ `cat output.txt`

```
"chars" 464  
"lines" 2  
"words" 78
```

Natural Join using MapReduce

Natural Join

- Joining database A and database B
 - Database A has attributes A1, A2, and A3
 - Database B has attributes A3, A4, and A5
- Find the records that agree on their attribute A3
 - i.e, the third attribute in database A and the first attribute in database B

Natural Join Example

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

R1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1

$R1 \bowtie S1 =$

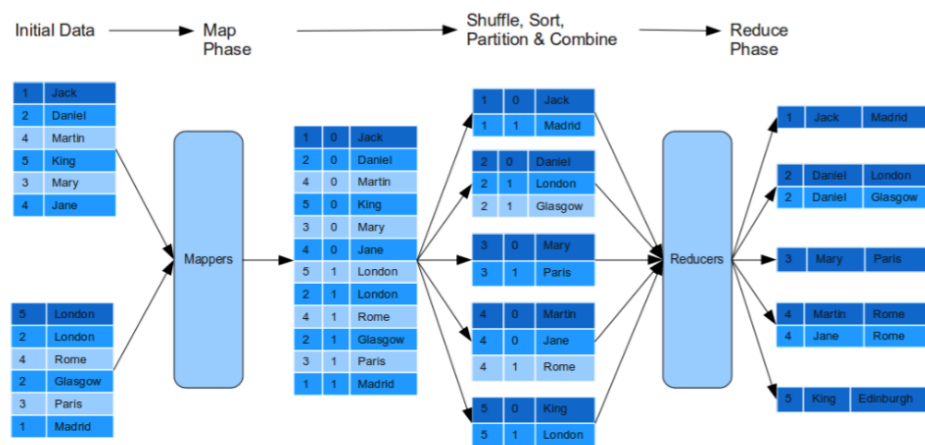
sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

Natural Join using MapReduce

- Mapper
 - For each record (A1, A2, A3) in database A
 - Produce a key-value pair (A3, ("A", (A1,A2)))
 - For each record (A3, A4, A5) in database B
 - Produce a key-value pair (A3, ("B", (A4,A5)))
- Reducer
 - For all of the pairs with same key
 - Construct all pairs comprising one with "A" and other with the first component "B"
 - e.g. , ("A", (A1,A2)) and ("B", (A4,A5))
 - Produce a record with (A1, A2, A3, A4, A5)

Natural Join using MapReduce

- Two-way joins (can be extended to multi-way joins)
 - The join operation actually happens on the reduce side



MapReduce Implementation

- Google has a proprietary implementation in C++
 - Bindings in Java, Python
- Hadoop is an open-source implementation in Java

Supporting Casts for Hadoop

- High Level languages for describing MapReduce applications
 - Pig (data exploration) - for data processing written in Pig Latin (dataflow language) developed by Yahoo
 - Hive (data warehousing) – Query language HQL developed by Facebook
 - MRJOB – supports Amazon's Elastic MapReduce service developed by Yelp
- HBase
 - Column-oriented distributed DBMS
- ZooKeeper
 - Coordination service for distributed applications